

000  
001  
002  
003  
004  
005  
006  
007  
008  
009  
010  
011  
012  
013  
014  
015  
016  
017  
018  
019  
020  
021  
022  
023  
024  
025  
026  
027  
028  
029  
030  
031  
032  
033  
034  
035  
036  
037  
038  
039  
040  
041  
042  
043  
044  
045  
046  
047  
048  
049  
050  
051  
052  
053

---

# Information Theoretic Estimators for Dependence in Time Series

---

Anonymous Author(s)  
Affiliation  
Address  
email

## Abstract

Recently, there has been much interest in designing and analyzing estimators for information theoretic functionals of probability density functions, under the assumption of independent and identically distributed (IID) data. However, estimators designed for IID data fail to capture relevant information when presented with time series data. We first present information theoretic functionals which take into account temporal dependence patterns, and then discuss methods for estimating these functionals. We then present an empirical evaluation of these estimators, in which we study the abilities of different estimators to detect dependencies in different types of synthetic data. This reveals novel patterns in the performance of these estimators over varying inputs, resulting in new understanding of the regimes in which different methods should be used to measure dependence between variables.

## 1 Introduction

Information theoretic functionals of probability density functions, including entropy, mutual information, divergence, and their conditional variants, are important nonparametric measures of variation, dependence, and distance between random variables. Over the past few years, much work has gone into designing and characterizing estimators for such functionals, using independent and identically distributed (IID) samples from each of the relevant variables.

However, many data sources, such as neurophysiological, economic, and meteorological data, consist of time series, where samples of each variable are not independent. When analyzing such data, it may often be useful to perform analyses using estimators for information theoretic functionals that were designed for IID data.

One particularly common use for information theoretic functionals is to estimate the strength of dependence between variables. If we simply treat two time series signals as sequences of IID samples and estimate the mutual information between them using previous methods, we are likely to fail to capture dependencies that manifest over multiple time steps. For example, to study functional connectivity in fMRI data, one might try to learn a graphical model over activity in different brain regions via the Chow-Liu or PC algorithm, which can use estimates of (conditional) mutual information as subroutines. If, in response to a stimulus, one brain region consistently activates a few seconds a second brain region, these brain regions ought to be considered as functionally connected, but typical information theoretic measures will fail to capture this dependence.

Thus, the broad goal of this work is to understand how methods using these functionals and their estimators can be adapted for time series data. In this paper, we focus on information theoretic functionals used to measure dependence, although some ideas of our work extend naturally to other functionals of interest. In particular, the two functionals we study in detail are transfer entropy and the mutual information rate, discussed in Sections 4 and 5, respectively.

## 2 Problem Setup and Notation

Suppose we observe two finite time series  $\mathcal{X} = \{X_i\}_{i=1}^n$  and  $\mathcal{Y} = \{Y_i\}_{i=1}^n$  (which can be written as a joint time series  $(\mathcal{X}, \mathcal{Y}) = \{(X_i, Y_i)\}_{i=1}^n$ ). We assume that the observations come from a stationary time series (i.e., that their joint densities are invariant to time shifts). In general, the  $2n$  variables in  $\{(X_i, Y_i)\}_{i=1}^n$  can have any dependence structure. However, in order to estimate quantities of interest, we will have to assume some independence in the form of Markov or mixing assumptions. For  $\beta \geq 1$ , a  $\beta$ -Markov assumption holds if, for each  $i \in \{\beta + 1, \dots, n\}$ ,

$$(X, Y)_i \perp\!\!\!\perp \{(X_j, Y_j)\}_{j=1}^{i-(\beta+1)} | \{(X_j, Y_j)\}_{j=i-\beta}^{i-1}, \quad (1)$$

$$(X, Y)_i \perp\!\!\!\perp \{X_j\}_{j=1}^{i-(\beta+1)} | \{X_j\}_{j=i-\beta}^{i-1} \quad \text{and} \quad (X, Y)_i \perp\!\!\!\perp \{Y_j\}_{j=1}^{i-(\beta+1)} | \{Y_j\}_{j=i-\beta}^{i-1}. \quad (2)$$

Mixing assumptions come in several forms and are generally more complicated, both to define and to verify for real data. Thus, we discuss them only where necessary in our theoretical work in the Supplementary Material, where they determine the rate of convergence of our estimators.

## 3 Related Work

Amblard and Michel (2012) review the histories of and theoretical relationships between direction information, Granger causality, and the mutual information rate. However, they do not give empirical evaluation of these quantities, and do not discuss estimation of the relevant information theoretic quantities. Otherwise, there has been much recent work on estimating information theoretic functionals in the case of IID data (Pál et al., 2010; Nguyen et al., 2010; Singh and Poczos, 2014; Krishnamurthy et al., 2014; Moon and Hero, 2014; Gao et al., 2014).

## 4 Transfer Entropy

One information theoretic functional that has been studied previously is the notion of transfer entropy (actually a case of conditional mutual information), defined for general time series  $\mathcal{X}$  and  $\mathcal{Y}$  as

$$T_{\mathcal{X} \rightarrow \mathcal{Y}}^n = H(Y_n | \{Y_i\}_{i=1}^{n-1}) - H(Y_n | \{X_i\}_{i=1}^{n-1}, \{Y_i\}_{i=1}^{n-1}) = I(Y_n; \{X_i\}_{i=1}^{n-1} | \{Y_i\}_{i=1}^{n-1}).$$

Thus, transfer entropy can be equivalently thought of as the reduction in uncertainty of  $Y_n$  upon learning  $\{X_i\}_{i=1}^n$  when one already knows  $\{Y_i\}_{i=1}^n$ , or the mutual information between the last point of  $Y$  and the previous values of  $X$  given the previous values of  $Y$ . Thus, it roughly measures the predictive power of  $X$  on  $Y$  after accounting for autocorrelation and can be used in this sense as a dependence measure among time series.

### 4.1 Estimating Transfer Entropy

The transfer entropy  $T_{\mathcal{X} \rightarrow \mathcal{Y}}^n$  of  $\{X_1\}_{i=1}^n$  on  $\{Y_1\}_{i=1}^n$  is a mutual information between the 1-dimensional variable  $Y_n$  and the  $(n-1)$ -dimensional variable  $X_{1:(n-1)}$  conditioned on the  $(n-1)$ -dimensional variable  $Y_{1:(n-1)}$  (a  $2n-1$  dimensional conditional mutual information). Without making parametric or independence assumptions, this makes estimating transfer entropy, very difficult for all but very small  $n$  (even when given several samples of the entire time series). To make estimation more tractable, we reduce the dimension of the problem by making a  $\beta$ -Markov assumption. Transfer entropy then reduces to

$$\begin{aligned} T_{\mathcal{X} \rightarrow \mathcal{Y}}^n &= H(Y_n | \{Y_i\}_{i=1}^{n-1}) - H(Y_n | \{X_i\}_{i=1}^{n-1}, \{Y_i\}_{i=1}^{n-1}) \\ &= H(Y_n | \{Y_i\}_{i=n-\beta}^{n-1}) - H(Y_n | \{X_i\}_{i=i-\beta}^{i-1}, \{Y_i\}_{i=i-\beta}^{i-1}) \\ &= I(Y_n; \{X_i\}_{i=i-\beta}^{i-1} | \{Y_i\}_{i=i-\beta}^{i-1}). \end{aligned}$$

which now involves at most  $\beta$ -dimensional variables (a  $2\beta + 1$  dimensional conditional mutual information). When  $\beta \ll n$ , this makes estimation tractable. In particular, estimating transfer entropy reduces to estimating a low-dimensional conditional entropy. This conditional entropy was then estimated using previously derived methods, as discussed in the Supplementary Material.

## 5 Mutual Information Rate

Yet another information theoretic measure of dependence between two time series is the mutual information rate (MIR), defined for a joint time series  $(\mathcal{X}, \mathcal{Y}) = \{(X_i, Y_i)\}_{i=1}^{\infty}$  by

$$I_R(\mathcal{X}, \mathcal{Y}) = \lim_{n \rightarrow \infty} \frac{I(X_1, \dots, X_n; Y_1, \dots, Y_n)}{n},$$

when this limit exists. Like transfer entropy, high dimensionality makes estimation difficult, except for short time series of many samples, and so, to make estimation tractable, we will assume  $\beta$ -Markov conditions (1) and (2) as above. We first discuss some results for the simpler case of estimating the entropy rate of a single time series, and then build up to estimating the MIR.

### 5.1 Estimating the Entropy Rate

Under  $\beta$ -Markov assumptions, then entropy rate  $H_R(\mathcal{X})$  reduces to

$$H_R(\mathcal{X}) = \lim_{n \rightarrow \infty} H(X_n | \{X_i\}_{i=1}^{n-1}) = \lim_{n \rightarrow \infty} H(X_n | \{X_i\}_{i=n-\beta}^{n-1}). \quad (3)$$

Thus, if we can estimate the conditional distribution of  $X_n | \{X_i\}_{i=n-\beta}^{n-1}$ , then we can attempt to use a plugin estimator. For an ergodic Markov chain with finite state space  $S$  (e.g., if  $\beta = 1$ <sup>1</sup> and the conditional distribution  $\mathbb{P}[X_{i+1} = x_\ell | X_i = x_j]$  is bounded below), the usual maximum likelihood estimate<sup>2</sup>

$$\hat{\pi}_{x, \{y_i\}_{i=1}^{\beta}} = \frac{\sum_{i=\beta+1}^n 1_{\{X_i=x, \{X_i\}_{i=n-\beta}^{n-1}=\{y_i\}_{i=1}^{\beta}\}}}{\sum_{i=\beta+1}^n 1_{\{\{X_i\}_{i=n-\beta}^{n-1}=\{y_i\}_{i=1}^{\beta}\}}}$$

of  $\mathbb{P}[X_n = x | \{X_i\}_{i=n-\beta}^{n-1} = \{y_i\}_{i=1}^{\beta}]$  (for each combination  $x, y_1, \dots, y_\beta \in S$ ) used for IID data is consistent (and, in fact, asymptotically normal). Ciuperca and Girardin (2005) show that, in case of an ergodic Markov chain with finite state space, this plug-in estimator is strongly consistent and asymptotically normal as  $n \rightarrow \infty$ .

### 5.2 Estimating the MIR

One might conjecture the most obvious analogue of (3) for MIR:

$$I_R(\mathcal{X}, \mathcal{Y}) \stackrel{?}{=} \lim_{n \rightarrow \infty} I(X_n; Y_n | \{(X_i, Y_i)\}_{i=n-\beta}^{n-1}). \quad (4)$$

However, (4) does not hold due to the more complicated effects of conditioning on mutual information (than on entropy). In particular, we have

$$I_R(\mathcal{X}; \mathcal{Y}) = \lim_{n \rightarrow \infty} \frac{I(X_1, \dots, X_n; Y_1, \dots, Y_n)}{n} \quad (5)$$

$$= \lim_{n \rightarrow \infty} \frac{H(X_1, \dots, X_n) + H(Y_1, \dots, Y_n) - H(X_1, \dots, X_n, Y_1, \dots, Y_n)}{n} \quad (6)$$

$$= H_R(\mathcal{X}) + H_R(\mathcal{Y}) - H_R(\mathcal{X}, \mathcal{Y}) \quad (7)$$

$$= \lim_{n \rightarrow \infty} H(X_n | \{X_i\}_{i=n-\beta}^{n-1}) + H(Y_n | \{Y_i\}_{i=n-\beta}^{n-1}) - H(X_n, Y_n | \{(X_i, Y_i)\}_{i=n-\beta}^{n-1}) \quad (8)$$

$$= \lim_{n \rightarrow \infty} H(X_n | \{(X_i, Y_i)\}_{i=n-\beta}^{n-1}) + H(Y_n | \{(X_i, Y_i)\}_{i=n-\beta}^{n-1}) + T_{\mathcal{Y} \rightarrow \mathcal{X}}^n + T_{\mathcal{X} \rightarrow \mathcal{Y}}^n - H(X_n, Y_n | \{(X_i, Y_i)\}_{i=n-\beta}^{n-1}) \quad (9)$$

$$\geq \lim_{n \rightarrow \infty} H(X_n | \{(X_i, Y_i)\}_{i=n-\beta}^{n-1}) + H(Y_n | \{(X_i, Y_i)\}_{i=n-\beta}^{n-1}) - H(X_n, Y_n | \{(X_i, Y_i)\}_{i=n-\beta}^{n-1}) \quad (10)$$

$$= \lim_{n \rightarrow \infty} I(X_n; Y_n | \{(X_i, Y_i)\}_{i=n-\beta}^{n-1}) \quad (11)$$

<sup>1</sup>The condition  $\beta = 1$  is not restrictive here, because any ergodic  $\beta$ -order Markov chain is equivalent to an ergodic first-order Markov chain with  $\beta$ -dimensional states.  $\beta$  continues to play a significant role in the convergence rates of the estimators, however.

<sup>2</sup>For any event  $A$ ,  $1_A$  denotes the indicator function of  $A$ .

where the inequality in (10) is due to the fact that the transfer entropy is non-negative, with equality precisely when  $X_n \perp\!\!\!\perp \{Y_i\}_{i=n-\beta}^{n-1} | \{X_i\}_{i=n-\beta}^{n-1}$  and  $Y_n \perp\!\!\!\perp \{X_i\}_{i=n-\beta}^{n-1} | \{Y_i\}_{i=n-\beta}^{n-1}$  (i.e., when  $T_{\mathcal{Y} \rightarrow \mathcal{X}}^n = T_{\mathcal{X} \rightarrow \mathcal{Y}}^n = 0$ , or, equivalently, when information is exchanged between  $\mathcal{X}$  and  $\mathcal{Y}$  only during simultaneous observations, with not lag effect). Fortunately, we can use the estimators derived above for the entropy rate to estimate  $H_R(\mathcal{X})$ ,  $H_R(\mathcal{Y})$ , and  $H_R(\mathcal{X}, \mathcal{Y})$  in (7), giving the estimator

$$\hat{I}_R(\mathcal{X}, \mathcal{Y}) = \hat{H}_R(\mathcal{X}) + \hat{H}_R(\mathcal{Y}) - \hat{H}_R(\mathcal{X}, \mathcal{Y}).$$

### 5.3 Information Rate and Transfer Entropy

Notice that line (9) of the above included the transfer entropies  $T_{\mathcal{Y} \rightarrow \mathcal{X}}^n$  and  $T_{\mathcal{X} \rightarrow \mathcal{Y}}^n$ . In particular, we see that MIR decomposes into

$$I_R(\mathcal{X}; \mathcal{Y}) = \lim_{n \rightarrow \infty} T_{\mathcal{Y} \rightarrow \mathcal{X}}^n + T_{\mathcal{X} \rightarrow \mathcal{Y}}^n + I(X_n; Y_n | \{(X_i, Y_i)\}_{i=n-\beta}^{n-1}).$$

Intuitively, the average information shared between  $\mathcal{X}$  and  $\mathcal{Y}$  can be decomposed into the (directed) lag effects of  $Y$  on  $X$  ( $T_{\mathcal{Y} \rightarrow \mathcal{X}}^n$ ), the (directed) lag effects of  $X$  on  $Y$  ( $T_{\mathcal{X} \rightarrow \mathcal{Y}}^n$ ), and the (undirected) simultaneous shared information of  $X$  and  $Y$  after accounting for history ( $I(X_n; Y_n | \{(X_i, Y_i)\}_{i=n-\beta}^{n-1})$ ). Contrasted from transfer entropy, MIR is thus a coarser, more holistic (undirected) measure of shared information between two time series. This will be further explored empirically in Section 6.

## 6 Empirical Results

We performed experiments on synthetic data to compare the behavior of estimators for transfer entropy, MIR, and the usual (IID) mutual information, as well as the behaviors of two approaches to estimating these quantities,  $k$ -nearest neighbor (KNN) estimators and those based on kernel density estimates (KDE estimators). Most previous empirical studies of information theoretic estimators (Singh and Póczos, 2014; Krishnamurthy et al., 2014; Pál et al., 2010; Moon and Hero, 2014; Gao et al., 2014) have compared the absolute errors of estimators for distributions where the true quantity can be computed analytically. In our experiments, we instead compare the power of using different estimators to detect dependencies at a fixed false positive rate. We do so for two reasons:

Firstly, we are interested in comparing the abilities of not only different estimators, but also different information theoretic quantities to capture dependence between variables. Transfer entropy, MIR, and mutual information measure fundamentally different quantities, and hence the errors of their estimators cannot be compared fairly.

Secondly, precise estimation of information theoretic functionals is (a) unnecessary for many practical applications and (b) subject to nuanced sources of error irrelevant to our evaluation. In most practical applications, the estimated mutual information between two variables is itself not interpretable or informative; rather, we are interested only how this quantity compares across different pairs of variables, through, for example, the ranking of the mutual information estimates over variables. Furthermore, for example, Gao et al. (2014) showed that classical KNN estimators tend severely underestimate the mutual information between variables which are strongly dependent.<sup>3</sup> On the other hand, the estimated mutual information still tends, on average, to monotonically increase with the true mutual information, and hence, given sufficient data, we can still reliably estimate the ranking of mutual information. As explained in the next section, our experiments rely only on the ranking of the mutual information estimates, and are thus more robust to issues of precise estimation such as that discussed in Gao et al. (2014). As a consequence of this robustness, we are also able to obtain good results with in our experiments with samples that are far smaller (by at least an order of magnitude) than previous experiments on estimating information theoretic quantities.

### 6.1 Permutation Tests for Independence

Suppose we observe  $n$  joint samples  $(\mathcal{X}_n, \mathcal{Y}_n) = \{(X_i, Y_i)\}_{i=1}^n$  from two time series  $\mathcal{X}$  and  $\mathcal{Y}$ . Suppose also that, for some quantity  $F(\mathcal{X}, \mathcal{Y})$ , we have a consistent estimator  $\hat{F}(\mathcal{X}_n, \mathcal{Y}_n)$  with

<sup>3</sup>Gao et al. (2014) also show that this can be corrected by performing local dimension reduction before applying the estimators. Intuitively, KDE estimators are also subject to this problem and can be likely corrected using locally adaptive kernels, but this has not yet been shown.

the property that, when  $\mathcal{X} \perp \mathcal{Y}$ , the distribution of  $\hat{F}(\mathcal{X}_n, \mathcal{Y}_n)$  is the same as the distribution of  $\hat{F}(\mathcal{X}_n, \mathcal{Y}_{n,\sigma})$ , where  $\mathcal{Y}_{n,\sigma} = \{Y_{\sigma(i)}\}_{i=1}^n$  for any permutation  $\sigma : \{1, \dots, n\} \rightarrow \{1, \dots, n\}$ . Under these conditions, it is easy to see that the following procedure is a valid  $\alpha$ -level hypothesis test for the null  $H_0 : \mathcal{X} \perp \mathcal{Y}$ :

1. For a large number  $P$ , generate  $P$  permutations  $\sigma_1, \dots, \sigma_P$  uniformly at random.
2. Compute the statistic  $\hat{p} := \frac{1}{P} \sum_{i=1}^P 1_{\{\hat{F}(X, Y) > \hat{F}(X, Y_{\sigma_i})\}}$
3. Reject  $H_0$  if and only if  $\hat{p} < \alpha$ .

That is, we reject  $H_0$  when  $\hat{F}(\mathcal{X}_n, \mathcal{Y}_n)$  is greater than  $1 - \alpha$  fraction of  $\hat{F}(\mathcal{X}_n, \mathcal{Y}_{n,\sigma})$ . If, additionally,  $F(\mathcal{X}, \mathcal{Y})$  satisfies  $F(\mathcal{X}, \mathcal{Y}) = 0$  whenever  $\mathcal{X} \perp \mathcal{Y}$  and  $F(\mathcal{X}, \mathcal{Y}) > 0$  otherwise, we expect this test to have some power against the alternative  $H_1$ , because permuting  $\mathcal{Y}$  randomly ought to break most of the dependence structure between  $\mathcal{X}$  and  $\mathcal{Y}$ , so that, if  $\mathcal{X} \not\perp \mathcal{Y}$ ,  $\hat{F}(\mathcal{X}_n, \mathcal{Y}_{n,\sigma}) \rightarrow 0$  while  $\hat{F}(\mathcal{X}_n, \mathcal{Y}_n) \rightarrow F(\mathcal{X}, \mathcal{Y}) > 0$ .

This type of permutation test is the basis of our empirical studies in this section, using  $T_{\mathcal{X} \rightarrow \mathcal{Y}}^n$ ,  $I$ , and  $I_R$  in place of  $F$ , and the respective estimators discussed above in place of  $\hat{F}$ .

## 6.2 Details of the Experiments

Each experiment used  $n = 300$  time samples, generated after 1000 identical samples that were discarded to ensure stationarity. Permutation tests used a type I error rate of  $\alpha = 0.05$  and  $P = 2000$  permutations. In all of the following experiment descriptions,  $\{\varepsilon_i\}_{i \in \mathbb{Z}}$  is IID standard Gaussian noise (i.e., each  $\varepsilon_i \sim \mathcal{N}(0, 1)$ ), and  $\{\delta_i\}_{i \in \mathbb{Z}}$  is IID zero-mean Gaussian noise with variance  $\sigma_Y$  (i.e., each  $\varepsilon_i \sim \mathcal{N}(0, \sigma_Y)$ ), for each  $\sigma_Y \in \{1, \dots, 10\}$  (as seen over the  $x$ -axes in Figure 1). Finally, for each experiment, we tried 3 different lags  $\beta \in \{1, 2, 3\}$ . The same  $\beta$  was used to generate the data and to estimate the transfer entropy and mutual information rate. Note that  $\beta$  does not affect on the mutual information between  $X_i$  and  $Y_i$  (except indirectly through the lagged dependence in Experiment 1).

**Experiment 1 (Simultaneous and Lagged Dependence):** In this experiment,  $\mathcal{X}$  and  $\mathcal{Y}$  depend on each other due to both simultaneous shared noise and a lagged direct dependence of  $\mathcal{Y}$  on  $\mathcal{X}$ . In particular, we set

$$X_{i+1} = \frac{X_i}{2} + \varepsilon_{i+1}$$

$$\text{and } Y_{i+1} = X_{i+1-\beta} + \varepsilon_{i+1} + \delta_{i+1}$$

Because  $X_{i-\beta}$  is predictive of  $Y_i$  and  $X_i$  shares noise  $\varepsilon_i$  with  $Y_i$ , we expect all three of transfer entropy, mutual information, and MIR to have some power here.

**Experiment 2 (Simultaneous Dependence):** In this experiment,  $\mathcal{X}$  and  $\mathcal{Y}$  depend on each other only through simultaneous shared noise. In particular, we set

$$X_{i+1} = \frac{X_i}{2} + \varepsilon_{i+1}$$

$$\text{and } Y_{i+1} = \varepsilon_{i+1} + \delta_{i+1}$$

Because  $X_{i-\beta}$  is never predictive of  $Y_i$  we expect transfer entropy to fail here. Because  $X_i$  shares noise  $\varepsilon_i$  with  $Y_i$ , we expect mutual information and MIR to have some power here.

**Experiment 3 (Lagged Dependence):** In this experiment,  $\mathcal{X}$  and  $\mathcal{Y}$  depend on each other only through a lagged direct dependence of  $\mathcal{Y}$  on  $\mathcal{X}$ . In particular, we set

$$X_{i+1} = \varepsilon_{i+1}$$

$$\text{and } Y_{i+1} = X_{i+1-\beta} + \delta_{i+1}$$

Because  $X_{i-\beta}$  is never predictive of  $Y_i$  we expect transfer entropy to fail here. Because  $X_i$  shares noise  $\varepsilon_i$  with  $Y_i$ , we expect mutual information and MIR to have some power here.

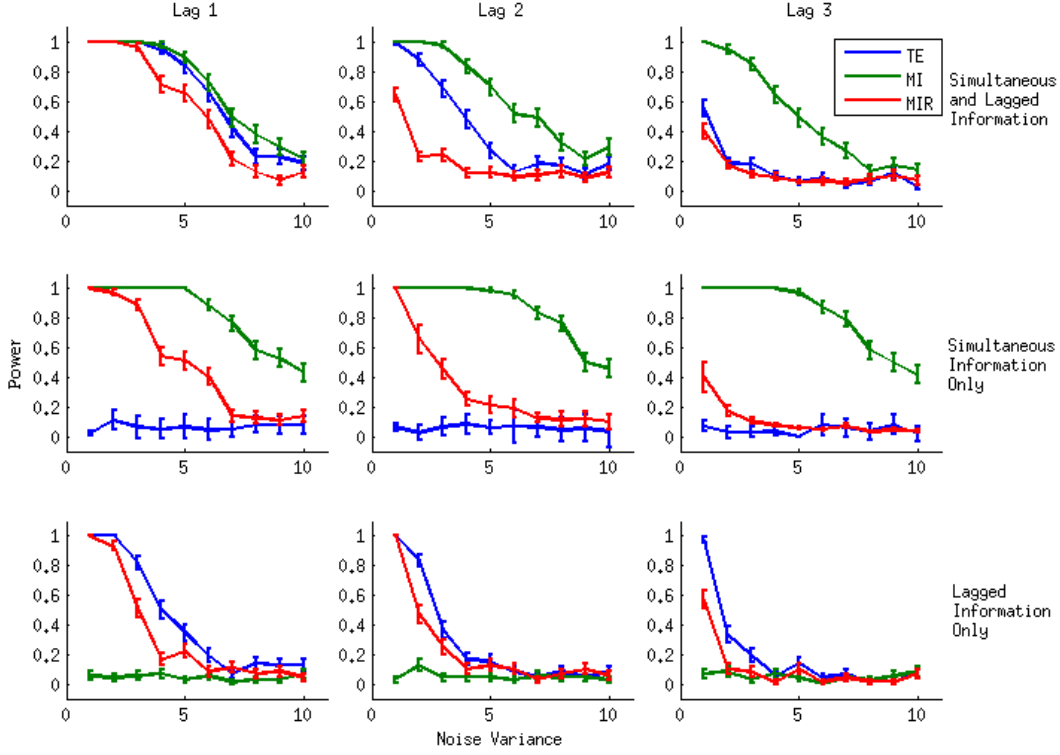


Figure 1: Graphical results of our experiments. Each plot shows the power of the permutation test for independence based on estimating each of transfer entropy, mutual information, and MIR (higher values indicate greater power), as a function of the noise in the dependence (i.e., the  $x$ -axes are roughly the noise-to-signal ratio). Plots to the right correspond to longer range dependencies within the time series (i.e., longer lag effects), and correspondingly weaker independence assumptions on the time series when computing  $\hat{T}_{\mathcal{X} \rightarrow \mathcal{Y}}^n$  and  $\hat{I}_R(\mathcal{X}, \mathcal{Y})$  (i.e., increasing  $\beta$ ). Plots in the same row have similar dependency structure (e.g., simultaneous or lagged dependence). Each curve plotted is averaged over 72 independent trials, and error bars show standard error over these trials.

### 6.3 Comparing measures of dependence

#### 6.3.1 Transfer Entropy, Mutual Information, and Mutual Information Rate Measure Different Dependence Structures

In several respects, the estimators behaved as anticipated by design. In Experiment 1, all three estimators were quite powerful (it is worth noting here that the noise in all of the above simulations is at least as, if not much more, powerful as the signal, so that demonstrating even moderate power is fairly notable). In Experiment 2, transfer entropy failed to detect dependence, while, in Experiment 3, mutual information failed to detect the dependence.

#### 6.3.2 MIR is the most general, least powerful measure

In all three experiments, MIR detected dependence (at least for low noise), which is consistent with the decomposition in (9) of MIR into two lagged components and a simultaneous component. However, because MIR attempts to measure multiple possible dependencies simultaneously, it uniformly has less power than other measures which detect precisely the dependence in the data (e.g., if there the data exhibit no lagged dependence, then the lagged dependence terms in MIR reduce to noise, reducing the power of the permutation test). Thus, in general, MIR should be used when there is little *a priori* reason to suspect a particular dependence structure, but a more powerful test can be devised otherwise.

## 6.4 Comparing estimators of entropy

### 6.4.1 KNN estimators scale better with dimensions

For  $\beta = 1$ , across all conditions, the kernel-based entropy estimator resulted in an average 20% more power than did the KNN estimator. However, as the lag  $\beta$  increased, the performance of KDE estimators dropped sharply, due to the strong effect of the curse of dimensionality on KDEs. On the other hand, the performance of KNN estimators was fairly consistent as  $\beta$  increased. This phenomenon was also observed in Moon and Hero (2014) as the dimension of the data increases, and the performance of both KDE and KNN approaches would likely benefit from the optimal ensemble modifications discussed in therein.

### 6.5 KDE estimators are more noise-tolerant

On the other hand, the performance of KNN estimators degraded more quickly than that of KDE estimators as independent Gaussian noise was added to each  $Y_i$ . That is, KDE estimators appear to be more tolerant to noise interfering in the relationship between the two variables.

Notably, for small  $\beta$ , our estimators actually performed *better* in the presence of some noise ( $\sigma = 1$ ) than in the noiseless case ( $\sigma = 0$ ). This is likely a reflection of the phenomenon mentioned above and discussed in detail in Gao et al. (2014), where the mutual information of strongly dependent variables is difficult to estimate without using certain corrections.

## 6.6 Conclusions to draw from the experiments

The primary practical lessons to glean from the empirical results are that

1. MIR is highly reliable without prior knowledge about the dependence structure.
2. Transfer entropy and mutual information can have higher power if the structure of the dependence is appropriate.
3. Kernel-based estimators outperform KNN estimators in low dimensions (small  $\beta$ ) and can be more robust to noise, but scale poorly with  $\beta$ .
4. The number of samples needed to have moderate power is not particularly large, even though quite a few samples may be needed to have small absolute error in estimating the quantity of interest.
5.  $\beta$  should be kept as small as possible, and has a significant effect on the number of samples needed for power.

## 7 Summary and Future Work

In this paper, we

1. discussed the intuition behind and relationships between some quantities that measure dependence between time series.
2. designed estimators for these quantities.
3. empirically compared the abilities of these estimators to detect different dependence structures under noise.
4. Compared the abilities of different estimators (KNN, KDE plug-in, etc.) to distinguish dependent and independent time series via a permutation test.

Additionally, in the Supplementary Material, we discuss an approach to proving convergence of these estimators under certain mixing assumptions on the time series.

Here, we focused primarily on estimating information theoretic functionals which measure dependence between time series. However, estimating functionals of probability densities from time series data can have applications to many other problems in machine learning and data analysis. For instance, we may wish to perform machine learning tasks over entire times series (as opposed to

individual points in time series), such as classifying entire time series as ordinary or anomalous. Since many machine learning algorithms can operate on distances between samples (rather than the original data points), one way to perform such tasks is to estimate distances or divergences between the distributions of the time series, which are often much more informative than the original time series (see, e.g., Póczos et al. (2012)). The convergence bounds developed in the Supplementary Material extends quite naturally to estimating arbitrary smooth functionals of probability densities, and so this should be developed, along with appropriate empirical evaluation.

Another potential application of these estimators is to visualizing dependencies within high-dimensional time series data. Suppose that we jointly measure  $M$  time series, and are interested in understanding structure of these dependencies. The pairwise dependency strengths can be computed to give an  $M \times M$  matrix, which can be compared to a Gram matrix in that it represents a sort of similarity between variables. This matrix can then be used in combinations with standard methods such as kernel  $k$ -means or multidimensional scaling (MDS), which take as input a similarity matrix. As a concrete example, suppose we obtain measurements at each of ten weather stations over the course of a year. In order to understand how much geographic proximity determines the interdependence of weather phenomena across regions, one may wish to visualize the measurements taken at different weather stations in terms of their similarity. Our proposed method using MDS gives a way of performing this visualization within, say, 2-dimensions, where results can be compared to geographic coordinates after rotating and scaling.

Finally, many other tests of independence and conditional independence of variables have been developed by others, for both time series and IID data. These include, a mutual information based bootstrap test for independence of time-series described in Wu et al. (2009), several independence tests described in (Gretton and Györfi, 2010), and conditional dependence measures (noting that, in principle, any conditional dependence measure can play the role of conditional mutual information in transfer entropy), e.g., (Bergsma, 2010) and (Zhang et al., 2012). The estimators and resulting permutation tests described here should be compared empirically and theoretically to these alternatives.

## References

- Pierre-Olivier Amblard and Olivier JJ Michel. The relation between granger causality and directed information theory: a review. *Entropy*, 15(1):113–143, 2012.
- Wicher Bergsma. Nonparametric testing of conditional independence by means of the partial copula. *Available at SSRN 1702981*, 2010.
- Gabriela Ciuperca and Valerie Girardin. On the estimation of the entropy rate of finite Markov chains. In *International Symposium on Applied Stochastic Models and Data Analysis (ASMDA)*, 2005.
- Shuyang Gao, Greg Ver Steeg, and Aram Galstyan. Efficient estimation of mutual information for strongly dependent variables. *arXiv preprint arXiv:1411.2003*, 2014.
- Arthur Gretton and László Györfi. Consistent nonparametric tests of independence. *The Journal of Machine Learning Research*, 11:1391–1423, 2010.
- A. Krishnamurthy, K. Kandasamy, B. Poczos, and L. Wasserman. Nonparametric estimation of renyi divergence and friends. In *International Conference on Machine Learning (ICML)*, 2014.
- K.R. Moon and A.O. Hero. Ensemble estimation of multivariate f-divergence. In *Information Theory (ISIT), 2014 IEEE International Symposium on*, pages 356–360, June 2014. doi: 10.1109/ISIT.2014.6874854.
- XuanLong Nguyen, Martin J Wainwright, and Michael I Jordan. Estimating divergence functionals and the likelihood ratio by convex risk minimization. *Information Theory, IEEE Transactions on*, 56(11):5847–5861, 2010.
- Dávid Pál, Barnabás Póczos, and Csaba Szepesvári. Estimation of rényi entropy and mutual information based on generalized nearest-neighbor graphs. In *Advances in Neural Information Processing Systems*, pages 1849–1857, 2010.
- Barnabás Póczos, Liang Xiong, and Jeff Schneider. Nonparametric divergence estimation with applications to machine learning on distributions. *arXiv preprint arXiv:1202.3758*, 2012.
- S. Singh and B. Poczos. Exponential concentration of a density functional estimator. In *Neural Information Processing Systems (NIPS)*, 2014.
- Edmond HC Wu, LH Philip, and Wai Keung Li. A smoothed bootstrap test for independence based on mutual information. *Computational Statistics & Data Analysis*, 53(7):2524–2536, 2009.



432 Kun Zhang, Jonas Peters, Dominik Janzing, and Bernhard Schölkopf. Kernel-based conditional  
433 independence test and application in causal discovery. *arXiv preprint arXiv:1202.3775*, 2012.  
434  
435  
436  
437  
438  
439  
440  
441  
442  
443  
444  
445  
446  
447  
448  
449  
450  
451  
452  
453  
454  
455  
456  
457  
458  
459  
460  
461  
462  
463  
464  
465  
466  
467  
468  
469  
470  
471  
472  
473  
474  
475  
476  
477  
478  
479  
480  
481  
482  
483  
484  
485